

The Denodo Platform

Enabling Intelligent Data Delivery

Gaining timely, effective insights from data is crucial in today's fast-paced business environment. It's all about innovation, making smart decisions, and keeping up with market changes. Imagine if everyone in your company could instantly access and integrate data from across the organization. Unfortunately, data landscapes are constantly evolving, with increasing cloud workloads and data spread across more environments than ever. The rising demand for data—fueled by technologies like generative AI (GenAI) and the emergence of agentic AI systems that can reason, plan, and act autonomously—further challenges data teams to keep pace.

Organizations are leveraging data lakehouses as a cornerstone of their modern data strategies, benefiting from scalable storage, powerful compute, and a unified platform for analytics and AI workloads. However, not all data resides in the lakehouse—and not all use cases are best served by centralizing it. Real-time operational reporting, AI applications that depend on live context, and scenarios requiring strict data sovereignty, often fall outside the lakehouse's traditional strengths. To address these gaps, organizations need a new approach—one that adds a layer of intelligence above the lakehouse and other existing systems to deliver data in the right form for each use case, in real time, regardless of where it resides or how it's structured.

Data Self-Service



Landsbankinn streamlined its data processes and enhanced accessibility, empowering over 80% of targeted users, including those with limited data skills, to make informed decisions.

Data Foundation for Improved Customer Experience



LeasePlan improved customer satisfaction by creating new programs like predictive vehicle maintenance, incorporating both internal and external data sources.




The Logical Approach to Data Management

The logical approach is based on a unified data delivery platform that simplifies how everyone in your business accesses the various data systems. It's about stripping away the complexity and serving up the data in formats that make sense to the user, while accelerating delivery and adhering to data governance rules.

The Denodo Platform is the leading logical data management solution, enabling intelligent data delivery from distributed sources. It leverages a universal semantic layer to present information in user-friendly formats, providing real-time, organization-wide access. With built-in security and governance, it enables the timely delivery of high-quality data, empowering business users with seamless access to a single, trusted data source.

Benefits include:

- **Unified Data Access:** Breaking down data silos and offering a centralized virtual data abstraction layer for easy access to data, regardless of location or format
- **Data Tailored for Business Users:** Translating complex data into user-friendly formats, simplifying access, querying, and analysis without requiring technical skills or IT help
- **Simplified Data Discovery:** Enhancing visibility and accessibility with rich metadata, detailed lineage information, and AI-powered, intuitive search capabilities, to revolutionize data discovery and usage
- **Real-Time Insights:** Delivering real-time data, ensuring timely access for faster decision-making and a competitive edge
- **Accelerated AI Development:** Speeds up GenAI and agentic AI projects by providing secure, real-time data access through a unified, developer-friendly interface.
- **Enhanced Data Security and Compliance:** Ensuring integrity, security, and reliability using advanced features and third-party-tool integration





IT Infrastructure Modernization

Sunbelt Rentals accelerated access to diverse data sources, boosting productivity by 200% by retrieving data from diverse systems, including over 500 ERP-system tables in under three minutes, without engineering support.



Enhance Operational Efficiency, Agility, and Resilience

City Furniture enhanced operations and profitability by developing a single source of truth across sales, supply chain, merchandising, and operations, resulting in improved market timing, expanded market share, and increased profits.



Centralized Governance, Risk, and Compliance

Seacoast Bank improved security and compliance efforts by using data for business process automation and making information readily available to key managers, including risk managers.

Denodo Platform Differentiators

- **Semantic Unification:** Provides a unified semantic layer for consistent, governed access to real-time data across all sources, with over 200 connectors to 3rd-party systems.
- **Personalized Self-Service:** Powers data marketplaces and tailored data products, enabling secure, self-service access without IT bottlenecks.
- **Federated Governance:** Enforces fine-grained, centrally managed data policies across environments to support privacy, compliance, and AI readiness.
- **Real-Time Data Delivery:** Connects to live source systems on demand, delivering current, trusted data for operational and AI workloads—without constant replication.
- **Embedded MPP Data Lake Engine:** Provides data consumers with a highly scalable and performant SQL engine, based on Presto, with which to access data lake data, while simplifying data access through an intuitive GUI interface.
- **Smart Query Acceleration:** Optimizes multi-source query execution through intelligent data materialization, cost-based optimization, and an AI-powered recommendations engine.
- **Financial Operations (FinOps):** Provides visibility into operations, for more effective management of cloud infrastructure costs.
- **Broad Delivery Options:** Provides more than JDBC or ODBC support for SQL queries by business users – developers can deliver data via a wider variety of options, including REST and GraphQL APIs, Kafka, and JMS message queues.
- **Data Preparation Wizard:** Empowers users of all backgrounds, especially business users to easily customize datasets within the data marketplace, enabling quick adaptation of data products for various use cases.

The Denodo AI SDK: Accelerating GenAI and Agentic AI Development

As organizations embrace GenAI, their focus will shift from simple content generation to the management of intelligent, autonomous systems—also known as agentic AI. These advanced applications depend on fast, secure access to context-rich enterprise data. That’s why a strong data foundation is essential.


The Denodo Platform provides that foundation. As a logical data management solution, it connects to all your data, regardless of its location, and makes it available in real time, with built-in security, governance, and a unified semantic layer. This provides AI models with data that is accurate, trusted, and aligned with the business, so AI initiatives can scale more quickly and responsibly.

To help developers harness this data foundation, Denodo offers the AI SDK, a powerful toolkit that accelerates the development of GenAI and agentic AI applications. The Denodo AI SDK provides a unified, secure interface for accessing and preparing enterprise data, removing the need for complex pipelines or custom integration code. It automatically manages data embedding, orchestration, and advanced security policies, so teams can focus on building intelligent applications that deliver real business value.

How the Denodo AI SDK supports AI initiatives:

- Denodo Query RAG:** Empowers users to ask natural language questions and get fast, accurate answers to fact-based queries—like what happened, when, or by how much—using real-time, governed data. It’s ideal for replacing manual reports and dashboards with an intuitive, conversational self-service way to access information.
- Denodo DeepQuery:** Tackles complex, open-ended questions that require reasoning, synthesis, or explanation—like understanding why something happened or identifying key drivers across systems. Delivers explainable, multi-step insights grounded in real-time data.
- MCP Support:** Enables seamless integration with any MCP-compliant client, so organizations can build and orchestrate multi-agent AI systems using open, interoperable standards. With a trusted data foundation provided by Denodo, enterprises can easily enable natural language and self-service access across their data ecosystems.

By bridging the gap between enterprise data and AI models, the Denodo AI SDK empowers organizations to build GenAI applications that are accurate, secure, and scalable—from intelligent assistants and copilots to advanced decision automation and beyond. And because it abstracts the complexity of modern GenAI tools, any developer familiar with REST APIs can start building powerful AI-driven applications with ease.



Denodo Assistant: Driving Business Efficiency and Better Outcomes

Denodo Assistant boosts business productivity by making data access and decision-making faster and smarter. Powered by the Denodo Platform's semantic layer and its extensive metadata, this intelligent assistant uniquely leverages technical, business, and operational metadata to enhance AI-driven data management tasks. By tapping into a unified and trusted data access layer, Denodo Assistant helps your teams work more efficiently, reducing the time and effort needed to find, understand, and access data. This intelligent assistant empowers technical and non-technical users, helping them make better decisions with guided recommendations, intuitive data access, and simplified workflows.

How Denodo Assistant benefits your business:

- Streamlined Data Access:** Enables teams to easily find and access the right data with natural language queries, eliminating the need for technical skills or IT support.
- Data Product Recommendations:** It provides intelligent suggestions based on past searches, helping users quickly identify the most relevant data for their needs and leading to faster, more informed decisions.
- Optimized Data Preparation:** Simplifies data preparation with AI-driven recommendations, enabling non-technical users to easily tailor data to their needs.
- Enriches Data Descriptions:** This feature automates the addition of descriptions to data views and columns, enhancing understanding and making it easier for teams to use the data effectively.
- Automatic Summary Recommendations:** Leverages machine learning and active metadata to suggest summary caches, optimizing query performance and efficiency according to business rules.
- Query Wizard Recommendations:** Guides users through query creation, enabling all skill levels to generate accurate queries with ease.
- Intelligent Autocompletion in VQL Shell:** Provides context-aware query suggestions, reducing errors and speeding up query writing.
- LLM-Based Data Enrichment Function:** Enables users to automatically summarize text, extract information, analyze sentiment, remove sensitive data, and translate text using large language models (LLMs), available as a function call in VQL.
- AI-Generated Join Recommendations:** Suggests intelligent join conditions between views, even without predefined relationships, helping data engineers build semantic models faster and with less manual effort.
- Natural Language to VQL Conversion:** Turns plain-language queries into VQL syntax, enabling anyone to build queries without VQL skills.
- AI-Driven Tagging for Governance:** It provides recommendations for tagging data assets, streamlining classification and simplifying security policy application for consistent, efficient governance.





Agora has all the robust capabilities of the Denodo Platform and offers them as a fully managed, cloud-based solution, simplifying data management while maintaining the platform's simplicity and power. Now available on both the AWS and Azure Marketplaces, Agora provides seamless, flexible deployment options, making it easier for organizations to discover, procure, and deploy. By offloading infrastructure and operations management, IT and operational teams can focus on business objectives and leverage the full potential of their data with agility, scalability, and cost efficiency.



COMPARING THE VARIOUS SUBSCRIPTION TIERS

There are several flexible subscription options to choose from, designed to suit the needs of various projects, from small departmental projects to enterprise-wide digital transformations.

	 DENODO PROFESSIONAL Small, single-use-case projects within individual departments	 DENODO STANDARD Multiple use cases within individual departments	 DENODO ENTERPRISE Enterprise-wide deployment for multiple use cases, and groups and large data volumes	 DENODO ENTERPRISE PLUS Comprehensive automation, collaboration, and advanced security for enterprise-wide deployments
Number of Data Sources Supported	5	Unlimited	Unlimited	Unlimited
FinOps Logging and Integrated Dashboard	✓	✓	✓	✓
Available as SaaS		✓	✓	✓
Integration with Version Control Systems (VCS)		✓	✓	✓
Smart Query Acceleration using summaries			✓	✓
VQL procedures			✓	✓
Advanced Diagnostic & Monitoring Tool			✓	✓
Data Marketplace			✓	✓
Data Marketplace: Dataset collaboration through endorsements, warnings, and deprecation notes			✓	✓
Integration with external Massively Parallel Processing (MPP) engines like Impala, Spark, and others			✓	✓
Integration with 3rd party data modeling tools (ER/Studio, Erwin, etc.)			✓	✓
Embedded MPP data lake engine				✓
Denodo Assistant, which automates tasks for data engineers and consumers using AI				✓
AI SDK: Unified interface for fast, secure data access and streamlined GenAI and agentic AI development				✓
Global security policies				✓
Import data governance tags from external catalogs				✓

Capabilities

E Indicates the feature(s) is available in Enterprise.

E+ Indicates the feature(s) is available in Enterprise Plus

DATA SOURCES

Relational Databases

- Generic (JDBC)
- IBM DB2 8.2, 9, 10, 11, and higher; 9,10, and 11 for z/OS
- Denodo Virtual DataPort - 8.0, 9
- Apache Derby 10 and higher
- Informix 7, 12
- Microsoft SQL Server 2000, 2005, 2008, 2008 R2, 2012, 2014, 2016, 2017, 2019, and higher
- MySQL: 4, 5, 8, and higher
- Oracle: 8i, 9i, 10g, 11g, 12c, 12c In-Memory, E-Business Suite 12, 18c, 19c, and higher
- PostgreSQL 11, 12, 13, 14, 15, 16, 17
- Sybase ASE / SAP ASE 12, 15
- MS Access
- Microsoft Dataverse
- Huawei GaussDB

In-Memory Databases

- SAP HANA 1.0, 2.0
- Oracle TimesTen 11g
- Oracle 12c In-Memory

Parallel Databases and Appliances

- ClickHouse
- Exasol 7.1
- Greenplum 4.2
- Vertica 7, 9
- Netezza 4.6, 5.0, 6.0, 7.0, and higher
- Oracle Exadata
- ParAccel 8.0.2 (by using ParAccel 2.5.0.0 JDBC3g with SSL driver)
- SybaseIQ
- SQreamDB
- Teradata 12, 13, 14, 15, 16, 17
- Yellowbrick

Cloud Data Warehouse / RDBMS

- Alibaba ApsaraDB for OceanBase MySQL
- Alibaba ApsaraDB for OceanBase Oracle
- Alibaba ApsaraDB RDS for MySQL
- Alibaba ApsaraDB RDS for PostgreSQL
- Alibaba ApsaraDB RDS for Microsoft SQL Server
- Alibaba ApsaraDB PolarDB for MySQL
- Alibaba ApsaraDB PolarDB for PostgreSQL
- Alibaba ApsaraDB AnalyticDB for MySQL
- Alibaba ApsaraDB AnalyticDB for PostgreSQL
- Alibaba MaxCompute
- Amazon Redshift
- Amazon Athena
- Amazon Aurora (MySQL and PostgreSQL)
- Amazon DynamoDB
- Azure Cosmos DB
- Azure SQL Database
- Azure Synapse SQL (aka Azure SQL Data Warehouse)
- Delta Lake
- Google AlloyDB
- Google BigQuery
- GCP (Google Cloud Platform) SQL for MySQL; GCP

SQL for PostgreSQL

- Google Spanner (Includes support for the service Spanner Data Boost)
- MongoDB Atlas
- Snowflake

Big Data

- Hive 0.13.0 (Hive Server 2); 1.1.0 (Hive Server 2); 2.0.0 (Hive Server 2); 3.1.2 and higher (Hive Server 3); Hive for Cloudera 1.1.0; Hive for Hortonworks 1.2.1
- Impala 1.2.4., 2.3, 3.x kudu
- Spark SQL 1.5, 1.6, 2.x, 3.x and higher
- PrestoDB
- Trino
- Databricks

NoSQL

- MongoDB
- Cassandra 3.x

Multi-Dimensional Sources

- Multidimensional database (generic)
- Azure Analysis Services
- SAP BI 3.x and BW 7.x
- Mondrian: 3.x
- IBM Cognos TM1
- MS SQL Server Analysis Services
- Essbase

Data Lake Storage/Formats

- S3
- Azure Data Lake Storage Azure Blob Storage
- Google Cloud Storage
- Parquet/Avro
- Delta/Iceberg
- Huawei Object Storage Service (OBS)

Web Services

- SOAP
- REST (XML, RSS, ATOM, JSON)
- OData v2.0 and v4.0

Flat and Binary Files

- CSV, pipe-delimited, regular expression-parsed
- MS Excel xls 97-2003
- MS Excel xlsx 2007 or later
- MS Access
- XML
- JSON
- SAS Files (SAS7BDAT)
- All files can be local or in remote filesystems, through FTP/ SFTP/FTPS, and in clear, zipped and/or encrypted format.

Indexes and unstructured content

- CMS, file systems, text
- Elasticsearch

Cloud, SaaS, Web Sources

- Adobe Analytics
- AWS
- Google Analytics
- Google Sheets
- Facebook
- LinkedIn
- MS Azure Data Lake
- MS Sharepoint
- MS Dynamics 365 Business Central / Customer Engagement
- Marketo
- ServiceNow
- Salesforce
- Twitter
- Workday
- many more through configurable JSON and XML adapters

Active Directory as Source or Leveraging Security

- LDAP v3
- Microsoft Active Directory 2003, 2008

Streaming/Messaging systems

- Kafka
- MQSeries
- SonicMQ
- ActiveMQ
- Tibco EMS
- Other JMS compatible services

Semantic Repositories

- Semantic repositories in Triple Stores / RDF accessed through SPARQL endpoints.

Packaged Applications

- SAP ERP/ECC (BAPIS and tables)
- Oracle E-Business Suite 12
- Siebel
- SAS 9.4

Hierarchical Databases

- Adabas (SOA Gateway and Denodo's SOAP connector): 5, 6

Denodo SDK for Custom Connectors

PUBLISHING OPTIONS

- SQL Based access via JDBC, ODBC and ADO.NET
- Web Services
 - REST
 - OData
 - Open API (a.k.a Swagger)
 - GraphQL
 - SOAP
- OAuth, OAuth 2.0 (JWT)
- SAML
- SSL
- WS-Security
- JMS listeners for message queues
- Denodo Scheduler for batch process and lite ETL

DATA MARKETPLACE

Cataloging

- Web-UI for seamless data discovery and exploration for business users
- Description and documentation
- Customizable business properties for richer metadata management
- Business categories and tags Intelligent search with smart ranking of results

Governance

- Graphical lineage spanning Denodo and Consumer tools such as Tableau and PowerBI
- Extended visibility into reports and dashboards
- Integrated request management (access, changes, data quality issues, etc.)
- Endorsement of datasets, comments, warnings, etc.
- Usage statistics: who uses what data, when and how
- Data profiling information

Self-Service

- Automatically matches UI language to browser settings
- Last-mile data preparation wizards for customizing datasets by non-technical users
- Full-featured SQL shell facilitates the execution of complex queries
- Export to CSV, Excel, and Tableau Data Extracts
- Save personal queries for easy access
- Share queries or publish them as new data products

PERFORMANCE OPTIMIZATIONS

- Smart query acceleration for analytics  
 - Aggregate Aware Summaries
- Massively parallel processing (MPP) integration   for query acceleration and caching
- Support for Arrow Flight SQL, enabling high-performance data transfer
- Full and partial aggregation and join pushdown, even in federated views
- Support for alternative data sources
- On-the-fly data movement for optimization
 - Option to restricted data movement for sensitive data
- Cost-based optimization (data statistics, data source indexes, data source execution model and parameters, network transfer rates)
- Pushdown of selections/projections/joins/groupby operations also on federated views
- Multiple join strategies
- Simplifying partitioned unions (partition pruning)
- and many more

DATA MATERIALIZATION OPTIONS

- Multi-mode selective data materialization (i.e. cache): full, partial, incremental, or total refresh, event-based or scheduled, configured at the view level, incremental queries for SaaS sources
- Directed Acyclic Graph (DAG) cache management automatically analyzes table dependencies and sequences replication tasks for optimal performance
- AlloyDB for PostgreSQL Amazon Athena
- Supports schema evolution without data loss
- Amazon Aurora MySQL

- Amazon Aurora PostgreSQL
- Amazon Redshift
- Azure SQL
- Azure Synapse SQL (previously known as Azure SQL Data Warehouse)
- Clickhouse
- Databricks (with Apache Arrow support for accelerated pushdowns)
- Exasol
- Google BigQuery
- GCP Cloud SQL for MySQL
- GCP Cloud SQL for PostgreSQL
- Hive 2.0.0; Hive 3.1.2, and higher (HiveServer2)
- IBM DB2 (8, 9, 10, 11, and higher for LUW; 9,10,11 for z/ OS)
- Impala 2.3; 3.x Kudu
- Microsoft Fabric Data Warehouse
- MS SQL Server (2000, 2005, 2008, 2008R2, 2012, 2014, 2016, 2017, 2019, and higher)
- MySQL 4, 5, 8, and higher
- Netezza 6, 7, and higher
- Oracle 8i, 9i, 10g, 11g, 12c, 12c In-Memory, 18c, 19c, and higher
- Oracle TimesTen 11g
- PostgreSQL 9, 10, 11, 12, and higher
- Presto
- SAP HANA 1 and 2
- Snowflake
- Spark SQL 2.x, 3.x, and higher
- Teradata 12, 13, 14, 15, 16, and 17
- Sybase
- Trino 4xx
- Vertica 7 and 9
- Yellowbrick
- Configurable “generic” adapter for other databases with JDBC drivers

DATA PIPELINES

- Remote Tables (created through UI or stored procedure)
- Denodo Scheduler
- VQL stored procedures

EMBEDDED MPP

- MPP engine based on Presto to accelerate access to data lake
- Graphical introspection of object storage (S3, ABFS, GFS, HDFS, etc.)
- Support for Parquet, Delta, and Iceberg
- New advanced optimization techniques to federate data lake content with any other data source

THIRD-PARTY MPP OPTIONS

- Impala
- Presto
- Spark 1.5, 1.6, 2.x
- Databricks 2.x

DATA GOVERNANCE

- Data source refresh, change impact analysis, dependency tree, full data lineage
- Denodo Governance Bridge: integration with IBM Information Governance Catalog

- API to publish metadata and lineage information to data governance tools like Informatica EDC, Collibra, etc.

SECURITY

Data in Motion – secure channels

- Using SSL/TLS
- Client-to-Denodo and Denodo-to-source
- Available for all protocols (JDBC, ODBC, ADO.NET and WS)

Data at Rest - secure storage

- Cache: third party database. Can leverage its own encryption mechanism
- Swapping to disk: serialized temporarily stored in a configurable folder that can be encrypted by the OS

Encryption/Decryption

- Support for custom decryption for files and web services
- Transparent integration with RDBMs encryption
- Encrypted metadata import/export

User and Role Based Including Integration with AD/ LDAP

- Row and column level authorization
- Advanced customizable masking
- Custom policies for specific security constraints and integration with external policy servers

Global Policies

- Tag-based security policies
- Support for RBAC and ABAC
- Dynamic Access Control Policies
- Column and row restrictions, multiple masking options, deny execution

Authentication

- Native and LDAP/Active Directory based Support for Kerberos and Windows SSO
- Kerberos
- NTLM
- OAuth, OAuth 2.0 (JWT)
- SAML
- Two-factor authentication (through supported identity providers: Okta, Duo, etc.)
- SSL
- WS-Security
- Pass-through session credentials to leverage existing source privileges

ADVANCED SEMANTICS

- Global security policies
- View- and column-level classification tags
- Support for the importing of external data governance tags from Collibra and other external data catalogs

DATA MODELING

- Design Studio: Web-based development studio for data modeling
- Desktop version also available
- Bottom-up and top-down (through Interface Views)
- Integration with third-party modeling tool  
 - ER/Studio Data Architect
 - ERwin Data Modeler
 - IBM InfoSphere Data Architect

- SAP PowerDesigner
- Sparx Systems Enterprise Architect

DATA QUALITY

- Library of transformation, filter, and matching functions and quality rules for validating, cleansing, enriching, standardizing, matching, and merging data
 - Extensible through custom functions
- Integration with external data quality tools

MONITORING

- Denodo Diagnostic and Monitoring Tool  (DMT) integrated in the Solution Manager
- Extensible usage and metadata dashboards integrated in Apache Superset
- FinOps dashboard to monitor and understand key metrics associated with cloud costs like egress, query cost, etc.
- Detailed monitoring information is available in logs for integration with log management tools like Splunk, ELK, Cloudwatch, etc.
- Monitoring is also available via SNMP and JMX standards. Therefore interoperate with most leading systems management packages (e.g., HP OpenView, Nagios, Zenoss, Osmius, IBM Tivoli and Microsoft WinRM)

DENODO ASSISTANT

Data Marketplace

Natural language queries with GenAI

- Support for OpenAI, Azure OpenAI, AWS Bedrock, and customer LLMs
- Personalized dataset recommendations
- Smart SQL-fragment autocomplete based on previous activity

Design Studio

- Summary recommendations for smart query acceleration
- View and column description suggestions
- Business friendly column name suggestions
- Diagnose incorrect queries and explain queries
- AI-generated recommendations for join conditions

VQL Function

- Automatically summarize, extract, analyze sentiment, remove sensitive data, and translate text using LLMs.

OPERATIONS

- Solutions Manager to automate operations and promotions tasks
 - Centralized management and distribution of updates to clients
 - Centralized management of license keys
 - Define promotion revisions and their dependencies and deploy them to a production cluster with zero downtime
 - Centralized management of data source properties and logs
 - REST API for automation of tasks from DevOps tools (e.g. jenkins)
- Integrated Infrastructure Management for Cloud (AWS)

- Creation and management of clusters: define type of EC2 instances, number of EC2 instances, etc.
- Creation of load balancers and Auto Scaling groups.
- Installation and launch of Denodo servers.
- Update the Denodo version
- Enable SSL in the Denodo servers.
- Multi-user development with version control integration
- Ingratiation with source management system such as Git
- Workspaces for isolated development and testing
- Resource Manager to limit and allocate resources to each session, role, or user in a way that optimizes resources utilization for each application
- Change resources priority
- Enforce limited timeouts or limits on number of rows
- Add daily quotas per minute/day/month: e.g. only 50 queries per day

DEPLOYMENT PATTERNS


- On-premises, private cloud, public cloud
 - On-premises, private cloud, public cloud
 - Basic single server configuration
 - HA cluster with load balancing (Active-Passive and Active-Active)
 - Shared or distributed local cache
 - Geographically distributed server environments
 - Multiple Denodo instances, peer-to-peer or multilayered
 - Containerization support through Docker
- Public cloud
 - Denodo Platform for AWS
 - Denodo Platform for Azure
 - Denodo Platform for GCP
 - Denodo Platform for Alibaba Cloud
- Auto-scaling support both in AWS and Azure
- Agora - the Denodo Cloud Service

OPERATING SYSTEMS

- **Microsoft:** Windows Server 2022, Windows Server 2019, Windows Server 2016, Windows 11, and Windows 10
- **Linux:** Amazon Linux 2023, Ubuntu 20.04 LTS or later
- CentOS 9.x and 8.x, Red Hat Enterprise Linux (RHEL) 9.x and 8.x, Oracle Linux 9.x and 8.x, SUSE Linux Enterprise 15.x

MINIMUM HARDWARE REQUIREMENTS

- Processor: Intel Xeon quad-core or similar. High-load scenarios or cases with complex calculations may require 8 cores or more.
- Physical memory (RAM): 16 gigabytes of memory so the Denodo server can allocate a runtime heap space up to 8 gigabytes.
- Disk space: Minimum: 5 gigabytes, Recommended: 100 gigabytes. Denodo only needs around 1 GB of disk space. If the cache is installed on the same server, more disk space will be required.



Denodo is a leader in data management. The award-winning Denodo Platform is the leading logical data management platform for transforming data into trustworthy insights and outcomes for all data-related initiatives across the enterprise, including AI and self-service. Denodo's customers in all industries all over the world have delivered trusted AI-ready and business-ready data in a third of the time and with 10x better performance than with lakehouses and other mainstream data platforms alone.

Visit www.denodo.com | Email info@denodo.com | Discover community.denodo.com

